


# Parallel Processing with Map Reduce

Chapter 07

Instructor: Houssein Dhayne  
houssein.dhayne@ul.edu.lb



# Table of Contents

- Introduction
- Google Search and MapReduce
- MapReduce Overview
- Sample MapReduce Application: WordCount
- MapReduce Programming
- MapReduce Jobs Execution
- Shuffle and Sort in MapReduce
- Progress and Status Updates in MapReduce
- Hadoop Streaming
- Hive Language




# Introduction to Parallel Processing

- Parallel processing is a powerful approach to efficiently handle large amounts of data.
- The ability to distribute tasks across multiple computing devices allows for faster data processing.

# Need for Parallel Processing in Big Data

- The increasing volume of big data requires innovative solutions for quick and efficient processing.
- Parallel processing enables the simultaneous execution of tasks across multiple computing devices.
- This approach is essential for handling vast datasets within reasonable time frames, addressing the challenges posed by big data.
- In this context, we explore the significance of parallel processing in managing big data effectively.



# How Google Search Uses MapReduce

- Google Search employs the MapReduce algorithm for efficient and rapid search query processing.
- Search queries are broken into sub-queries using keywords, and a Map program is generated for each sub-query.
- These Map programs are distributed to data nodes globally, where they process data and return results, demonstrating the scalability and effectiveness of MapReduce in real-world applications.

# Overview of MapReduce


- MapReduce serves as a parallel programming framework for large-scale data processing.
- It operates with minimal data movement on distributed file systems like Hadoop clusters, enabling near real-time results.
- The prerequisite for MapReduce programming includes applications suitable for parallel processing and data expressible in key-value pairs.

## USES OF MAPREDUCE



# Prerequisites for MapReduce Programming

- Two major prerequisites for MapReduce programming are:
  - The application must lend itself to parallel programming.
  - The data for the application must be expressible in key-value pairs.
- Drawing parallels with UNIX sequences, the structure of MapReduce processing is introduced, emphasizing the division of tasks and the intermediate results.



# Key-Pair Data Structure in MapReduce

- MapReduce relies on a key-pair data structure, where data is represented as key-value pairs.
- This structure facilitates the distribution of tasks and results among computing devices.
- Understanding and leveraging the key-pair data structure is crucial for effective implementation of MapReduce programs.

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

# MapReduce Processing Structure - UNIX Sequence Comparison

- MapReduce processing can be likened to UNIX sequences (pipes) structure.
- A concrete example is illustrated with a UNIX command:
  - `grep | sort | count myfile.txt`
- This comparison aids in understanding the flow of MapReduce tasks, highlighting the Map, Sort, and Reduce steps.

Grep
We
are
going
to
a
picnic
near
our
house
Many
of
our
friends
are
coming
You

Sort
a
Are
Are
Are
Coming
Friends
Fun
Going
Have
House
Join
Many
Near
Of
Our
Picnic
To
Us

WordCount	
A	1
Are	3
Coming	1
Friends	1
Fun	1
Going	1
Have	1
House	1
Join	1
Many	1
Near	1
Of	1
Our	2
Picnic	1
To	2
Us	1


# Program Structure for MapReduce Programs

- MapReduce programs are structured with a Map program and a Reduce program.
- The Map program processes input data and produces intermediate key-value pairs.
- The Reduce program takes these intermediate results, performs further processing, and produces the final output.
- This structured approach ensures efficient parallel processing of large datasets.

# Key Characteristics of MapReduce Data Processing


Key characteristics of MapReduce data processing are:

- Map step distributes the full job into smaller tasks processed on separate computers.
- Intermediate results from the Map step are considered as key2-value2 pairs.
- The Reduce step reads and combines the intermediate results to produce the final output in key2-value3 format.



# Sample MapReduce Application - WordCount

- A practical example of a MapReduce application is presented: WordCount.
- The scenario involves identifying unique words and their frequency in a text.



We are going to a picnic near our house.

1

Many of our friends are coming.

1

You are welcome to join us.

1

We will have fun.

1

Key2	Value2
we	1
are	1
going	1
to	1
a	1
picnic	1
near	1
our	1
house	1

Key2	Value2
many	1
Of	1
Our	1
friends	1
Are	1
coming	1

Key2	Value2
you	1
Are	1
Welcome	1
To	1
Join	1
Us	1

Key2	Value2
we	1
will	1
have	1
fun	1

- The sort process inherent within MapReduce will sort each of the intermediate files, and produce the following sorted key2-value2 pairs:

<b>Key2</b>	<b>Value2</b>
a	1
Are	1
Going	1
House	1
Near	1
Our	1
Picnic	1
to	1
We	1

<b>Key2</b>	<b>Value2</b>
Are	1
coming	1
friends	1
many	1
Of	1
our	1

<b>Key2</b>	<b>Value2</b>
Are	1
Join	1
To	1
Us	1
welcome	1
you	1

<b>Key2</b>	<b>Value2</b>
fun	1
have	1
we	1
will	1

## Reduce

- The Reduce function will read the sorted intermediate files, and combine the counts for all the unique words, to produce the following output.
- The keys remain the same as in the intermediate results.
- However, the values change as counts from each of the intermediate files are added up for each key.
- For example, the count for the word ‘are’ goes up to 3.

Key2	Value3
a	1
are	3
coming	1
friends	1
fun	1
going	1
have	1
house	1
join	1
many	1
near	1
of	1
our	2
picnic	1
to	2
us	1

# MAPREDUCE PROGRAMMING



# MapReduce Data Types and Formats

- MapReduce operates with a simple data processing model using key-value pairs.
  - map:  $(K_1, V_1) \rightarrow \text{list } (K_2, V_2)$
  - reduce:  $(K_2, \text{list}(V_2)) \rightarrow \text{list } (K_2, V_3)$
- Map and reduce functions have specific input and output types.
- Data formats, such as flat text files and databases, can be processed by MapReduce.

## Writing MapReduce Programs -1-

- In Java, the Map function is represented by the generic Mapper class.
- It uses four parameters: (input key, input value, output key, output value).
- This class uses an abstract map ( ) method. This method receives the input key and input value.
- It would normally produce output key and output value.
- A Mapper commonly performs input format parsing, projection (selecting the relevant fields), and filtering (selecting the records of interest).
- The Reducer typically combines (adds or averages) those values.

## Writing MapReduce Programs -2-

```
// Mapper: Emits (word, 1) for each occurrence
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {
    public void map(...) {
        // Split input line into words and emit (word,1)
    }
}

// Reducer: Sums counts for each word
public class SumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(...) {
        int sum = 0;
        for (IntWritable value : values) { sum += value.get(); }
        context.write(key, new IntWritable(sum));
    }
}
```

# Writing MapReduce Programs - step-by-step logic

- The big document is split into many segments. The Map step is run on each segment of data. The output will be a set of (key, value) pairs. In this case, the key will be a word in the document.
- The system will gather the (key, value) pair outputs from all the mappers, and will sort them by key. The sorted list itself may then be split into a few segments.
- A Reducer task will read the sorted list and produce a combined list of word counts.

```
map(String key, String value):  
    for each word w in value:  
        EmitIntermediate(w, "1");  
  
reduce(String key, Iterator values):  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    EmitAsString(result));
```


# MapReduce Program Optimization - Using Combiners

- To optimize MapReduce programs, using combiners is introduced.
- Combiners are functions that operate on the output of the Map function before it reaches the Reducer.
- For example, a combiner is introduced with the `combine` method. The combiner performs a partial aggregation of counts for each word locally on each Mapper node before sending the data to the Reducer.
  - This helps reduce the amount of data transferred to the Reducer, improving overall efficiency. The `combine` method has the same logic as the `reduce` method in the Reducer, but it operates locally on each Mapper node.

# MAPREDUCE JOBS EXECUTION



# Synchronization of Tasks in MapReduce



# MapReduce Job Execution- Overview

- **Job Components:**
  - MapReduce job comprises the Map and Reduce programs along with input datasets.
- **Job Coordination:**
  - A master program (Job Tracker on the NameNode) continuously monitors job progress.
  - Job coordination ensures that the distributed tasks are executed efficiently, reliably, and in a balanced manner throughout the cluster.
- **Distributed Processing:**
  - Map and Reduce tasks are executed on DataNodes hosting data fragments.
- **Resource Management:**
  - Communication and resource allocation among nodes are managed by YARN.

# Job Tracker, YARN, and Data Locality

- **Job Tracker:**

- Resides on the NameNode; responsible for coordinating job execution.
- Assigns tasks based on the location of data to minimize network traffic.

- **YARN:**

- Acts as the resource manager; allocates containers and schedules tasks across the cluster.

- **Data Locality:**

- The system schedules tasks on nodes where the data already resides (e.g., if Node A holds data (x, y, z), Map tasks for that data run on Node A).

# Task Execution: Map, Shuffle, and Reduce

- **Input Splits:**

- The dataset is divided into fixed-size pieces called input splits.
- One map task is created per input split.

- **Map Phase:**

- Each map task applies the user-defined Map function to process records and produce intermediate key-value pairs.

- **Shuffle and Sort:**

- Intermediate data is buffered, sorted by key, and partitioned for the reduce phase.
- This sorting is necessary for efficient processing during the Reduce phase.
- An optional combiner can perform local aggregation to reduce data transfer.

- **Reduce Phase:**

- Reduce tasks merge sorted outputs from map tasks to produce the final results.

# Job and Task Management

- **Task Tracker:**

- Each DataNode runs a Task Tracker to monitor assigned Map and Reduce tasks.
- On task completion, the Task Tracker notifies the Job Tracker.

- **Job Scheduling:**

- YARN's scheduler dynamically allocates resources for each task based on current cluster capacity.

# Managing Failures and Ensuring Resilience

- **Task Failures:**

- If a task fails (e.g., due to a runtime error), the Task Tracker reports the error to the Job Tracker, which then reschedules the task on a different node.

- **Application Master Failures:**

- The entire MapReduce job may be restarted if the Application Master (running on YARN) fails, up to a configurable limit.

- **Node Failures:**

- YARN detects DataNode failures via missing heartbeats and reallocates tasks from the failed nodes to healthy ones.

**These mechanisms ensure continuous processing even in the presence of hardware or software failures.**

# Testing MapReduce Jar: Word Count Example

Input File: wordscount.txt

Output Directory: wordscountout

Jar File: WordCount.jar

Main Class: org.WordCount.WordCountDriver

- Step 1: Upload the input file (wordscount.txt) to the Hadoop Distributed File System (HDFS).  
`hdfs dfs -put "/usr/mybigdata/wordscount.txt" /user/test/`
- Execute the MapReduce application by specifying the input file (wordscount.txt) and the output directory (wordscountout).  
`hadoop jar WordCount.jar org.WordCount.WordCountDriver /user/test/wordscount.txt /user/test/wordscountout`


# Hive Overview

Data warehouse infrastructure tool  
to process structured data in  
Hadoop.



# Introduction to Apache Hive

- Apache Hive is a data warehouse and ETL (Extract, Transform, Load) tool.
- Provides an SQL-like interface for users to interact with Hadoop Distributed File System (HDFS).
- Built on top of the Hadoop ecosystem, facilitating data query and analysis.
- **Apache Hive** makes querying and managing large datasets easier through SQL-like syntax called HiveQL. Hive translates these queries into **MapReduce jobs** under the hood to process data.



## Development and Contributors

- Initially developed by Facebook and later adopted by Amazon, Netflix, and others.
- Delivers standard SQL functionality for analytics.
- Not suitable for Online Transactional Processing (OLTP) workloads.
- Commonly used for data warehousing, and analyzing large datasets.
- Facilitates reading, writing, and handling wide datasets stored in distributed storage.

# Modes of Hive

- **Local Mode:**

- Used in pseudo mode with a single data node.
- Suitable for smaller datasets on a local machine.

- **MapReduce Mode:**

- Used with multiple data nodes in a Hadoop cluster.
- Executes queries in parallel for enhanced performance on large datasets.

## Characteristics of Hive

- Databases and tables must be created before loading data.
- Hive as data warehouse is built to manage and query only structured data which is residing under tables.
- Uses directory structures to improve performance on certain queries.
- Compatible with various file formats (TEXTFILE, SEQUENCEFILE, ORC, RCFILE).

## Features of Hive

- It provides indexes, including bitmap indexes to accelerate the queries.
- Built in user-defined functions (UDFs) to manipulation of strings, dates, and other data-mining tools.
- It stores schemas in a database and processes the data into the Hadoop File Distributed File System (HDFS).
  - Schema in Database: Hive stores the metadata (like table names, column types, partitions, and locations) in a Metastore — usually backed by a database like MySQL or PostgreSQL.
  - Data in HDFS: The actual data is stored in HDFS (Hadoop Distributed File System), while Hive queries act as a layer on top to process that data.

# Advantages of Apache Hive

- **Scalability:**
  - Designed to handle large volumes of data.
- **SQL-Like Interface:**
  - HiveQL is similar to SQL, making it familiar for SQL users.
- **Integration with Hadoop Ecosystem:**
  - Seamless integration with other Hadoop tools like Pig, MapReduce, and Spark.
- **Supports Partitioning and Bucketing:**
  - Improves query performance by limiting data scanned.

# Disadvantages of Apache Hive

- **Limited Real-time Processing:**
  - Designed for batch processing, not real-time data processing.
- **Slow Performance:**
  - Can be slower compared to traditional relational databases.
- **Steep Learning Curve:**
  - Requires knowledge of Hadoop and distributed computing.
- **Lack of Transaction Support:**
  - Does not support transactions, impacting data consistency.
- **Limited Flexibility:**
  - Specific to Hadoop, limiting usability in other environments.

## Hive Language Capabilities

- Hive's SQL provides standard SQL operations such as SELECT, FROM, WHERE, JOIN, GROUP BY, and ORDER BY.
- Results can be stored in tables or HDFS files, offering flexibility.
- Comparisons between Hive and traditional relational databases highlight its advantages for specific use cases.

# Creating a Database and Table

- **Create a database named 'school'**

```
CREATE DATABASE IF NOT EXISTS school;
```

- **Switch to the 'school' database**

```
USE school;
```

- **Create a table named 'student' within the 'school' database**

```
CREATE TABLE student (
    student_id INT,
    name STRING,
    age INT,
    grade STRING
);
```

# Inserting Data into a Table

- Insert data into the 'student' table

```
INSERT INTO student VALUES  
(1, 'John Doe', 20, 'A'),  
(2, 'Jane Smith', 22, 'B'),  
(3, 'Bob Johnson', 21, 'C');
```

- Overwrite existing data in the 'student' table

```
INSERT OVERWRITE TABLE student VALUES  
(1, 'Alice Johnson', 23, 'A'),  
(4, 'Charlie Brown', 24, 'B'),  
(5, 'Eva White', 22, 'A');
```

# **INSERT OVERWRITE with Immutable Files and Blocks in Hadoop**

## **How INSERT OVERWRITE Works with Immutable Files and Blocks in Hadoop?**

- INSERT OVERWRITE involves overwriting existing data in Hive.
- Hadoop follows an immutable data model - once written, data is not modified in-place.
- New set of data files is created, and the old ones are not directly modified.

# Loading Data into Hive

The following Hive script creates a student table to store structured data in a text file format.

The table uses a comma (',') as a delimiter to separate fields, making it easy to load and query data efficiently.

- Drop table student;
- CREATE TABLE student (
- student\_id INT,
- name STRING,
- age INT,
- grade STRING )
- ROW FORMAT DELIMITED
- FIELDS TERMINATED BY ','
- STORED AS TEXTFILE;

## Loading Data into Hive

- To load a file into Hive, you typically use the LOAD DATA statement.
- This statement is used to load data from a local file or HDFS (Hadoop Distributed File System) into a Hive table.

-- Load data into the Student' table from a local file

```
LOAD DATA LOCAL INPATH '/path/to/data.txt' INTO TABLE Student;
```

-- Load data into the Student table from HDFS

```
LOAD DATA INPATH '/hdfs/path/to/data.txt' INTO TABLE student;
```

## Creating and Loading Data into the wordcount Table in Hive

- CREATE TABLE wordcount (
    - word STRING,
    - count INT
  - )
  - ROW FORMAT DELIMITED
  - FIELDS TERMINATED BY ''
  - STORED AS TEXTFILE;
- 
- LOAD DATA INPATH '/user/test/wordscountout/part.csv' INTO TABLE wordcount;

# Conclusion

- MapReduce stands out as the leading parallel processing framework for Big Data applications.
- Its effectiveness lies in its ability to process large, divisible datasets, represented in `<key, value>` pair format.
- High-level languages like Hive and Pig enhance the ease of MapReduce programming.
- The presentation has covered fundamental MapReduce concepts, testing strategies, optimization techniques, and complementary tools like Hadoop Streaming and Hive.